

EDA & ML Models-(LinearRehgression, LogisticRegression, Kmeans, RandomForest)

January 12, 2026

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn.cluster import KMeans
from sklearn.decomposition import PCA

from sklearn.ensemble import RandomForestClassifier, RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error,
    classification_report, confusion_matrix, accuracy_score, roc_auc_score,
    roc_curve
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, OneHotEncoder, StandardScaler

from sklearn.impute import SimpleImputer
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
```

```
[2]: df = pd.read_csv("ncr_ride_bookings.csv")
df.head()
```

```
[2]:      Date        Time   Booking ID Booking Status Customer ID \
0  2024-03-23  12:29:38 "CNR5884300" No Driver Found "CID1982111"
1  2024-11-29  18:01:39 "CNR1326809"     Incomplete "CID4604802"
2  2024-08-23  08:56:10 "CNR8494506"     Completed "CID9202816"
3  2024-10-21  17:17:25 "CNR8906825"     Completed "CID2610914"
4  2024-09-16  22:08:00 "CNR1950162"     Completed "CID9933542"

      Vehicle Type       Pickup Location      Drop Location Avg VTAT Avg CTAT \
0             eBike           Palam Vihar            Jhilmil      NaN      NaN
1          Go Sedan         Shastri Nagar      Gurgaon Sector 56      4.9     14.0
2              Auto            Khandsa        Malviya Nagar      13.4     25.8
3  Premier Sedan  Central Secretariat        Inderlok      13.1     28.5
```

```

4           Bike      Ghitorni Village      Khan Market      5.3      19.6
... Reason for cancelling by Customer Cancelled Rides by Driver \
0 ...                   NaN                  NaN
1 ...                   NaN                  NaN
2 ...                   NaN                  NaN
3 ...                   NaN                  NaN
4 ...                   NaN                  NaN

Driver Cancellation Reason Incomplete Rides  Incomplete Rides Reason \
0                   NaN                  NaN                  NaN
1                   NaN                  1.0      Vehicle Breakdown
2                   NaN                  NaN                  NaN
3                   NaN                  NaN                  NaN
4                   NaN                  NaN                  NaN

Booking Value Ride Distance Driver Ratings Customer Rating \
0          NaN        NaN          NaN          NaN
1       237.0      5.73          NaN          NaN
2       627.0     13.58          4.9        4.9
3       416.0     34.02          4.6        5.0
4       737.0     48.21          4.1        4.3

Payment Method
0          NaN
1          UPI
2      Debit Card
3          UPI
4          UPI

```

[5 rows x 21 columns]

[3]: df.shape

[3]: (150000, 21)

[4]: df.dtypes

| | |
|-----------------|---------|
| Date | object |
| Time | object |
| Booking ID | object |
| Booking Status | object |
| Customer ID | object |
| Vehicle Type | object |
| Pickup Location | object |
| Drop Location | object |
| Avg VTAT | float64 |

```
Avg CTAT                      float64
Cancelled Rides by Customer    float64
Reason for cancelling by Customer   object
Cancelled Rides by Driver      float64
Driver Cancellation Reason    object
Incomplete Rides              float64
Incomplete Rides Reason        object
Booking Value                 float64
Ride Distance                 float64
Driver Ratings                 float64
Customer Rating                float64
Payment Method                  object
dtype: object
```

```
[5]: df.isnull().sum()
```

```
Date                      0
Time                      0
Booking ID                 0
Booking Status              0
Customer ID                 0
Vehicle Type                0
Pickup Location              0
Drop Location                0
Avg VTAT                   10500
Avg CTAT                   48000
Cancelled Rides by Customer 139500
Reason for cancelling by Customer 139500
Cancelled Rides by Driver   123000
Driver Cancellation Reason 123000
Incomplete Rides             141000
Incomplete Rides Reason     141000
Booking Value                48000
Ride Distance                48000
Driver Ratings                57000
Customer Rating               57000
Payment Method                 48000
dtype: int64
```

```
[6]: df = df[['Booking Status', 'Vehicle Type', 'Pickup Location', 'Drop Location',  
           ↴'Avg VTAT', 'Avg CTAT', 'Booking Value', 'Ride Distance', 'Driver Ratings',  
           ↴'Customer Rating', 'Payment Method']]
```

```
[7]: df.fillna(df.median(numeric_only=True), inplace=True)
```

```
[8]: df.isnull().sum()
```

```
[8]: Booking Status      0  
Vehicle Type          0  
Pickup Location       0  
Drop Location         0  
Avg VTAT              0  
Avg CTAT              0  
Booking Value         0  
Ride Distance         0  
Driver Ratings        0  
Customer Rating       0  
Payment Method        48000  
dtype: int64
```

```
[9]: df.duplicated().sum()
```

```
[9]: np.int64(320)
```

```
[10]: df.describe()
```

```
[10]:          Avg VTAT      Avg CTAT  Booking Value  Ride Distance \\\n  count  150000.000000  150000.000000  150000.000000  150000.000000  
  mean     8.445407    29.037753   478.121220    24.343568  
  std      3.639311    7.343053   329.339976   11.554362  
  min      2.000000    10.000000   50.000000    1.000000  
  25%     5.600000    25.000000   319.750000   17.360000  
  50%     8.300000    28.800000   414.000000   23.720000  
  75%    11.000000    32.900000   521.000000   30.650000  
  max     20.000000    45.000000  4277.000000   50.000000  
  
          Driver Ratings  Customer Rating  
  count  150000.000000  150000.000000  
  mean     4.257215    4.440842  
  std      0.345619    0.347835  
  min      3.000000    3.000000  
  25%     4.200000    4.300000  
  50%     4.300000    4.500000  
  75%     4.300000    4.600000  
  max     5.000000    5.000000
```

```
[11]: num_cols = df.select_dtypes(include=[np.number]).columns.tolist()  
cat_cols = df.select_dtypes(include=["object", "category", "bool"]).columns.  
           tolist()
```

```
[12]: num_cols
```

```
[12]: ['Avg VTAT',  
      'Avg CTAT',
```


```
'Booking Value',
'Ride Distance',
'Driver Ratings',
'Customer Rating']
```


```
[13]: cat_cols
```

```
[13]: ['Booking Status',
'Vehicle Type',
'Pickup Location',
'Drop Location',
'Payment Method']
```


Univariate Analysis -----


```
[14]: for col in num_cols[:5]:
    plt.figure()
    sns.histplot(df[col].dropna(), bins=30, kde=True)
    plt.title(f"Histogram of {col}")
    plt.show()
```






Histogram of Booking Value









```
[15]: for col in cat_cols[:5]:
    plt.figure()
    df[col].value_counts().head(10).plot(kind="bar")
    plt.title(f"Top Categories of {col}")
    plt.show()
```


Top Categories of Booking Status




Top Categories of Vehicle Type




Top Categories of Pickup Location



Top Categories of Drop Location




Top Categories of Payment Method



Correlation -----

```
[16]: if len(num_cols) > 1:  
    plt.figure(figsize=(8,6))  
    sns.heatmap(df[num_cols].corr(), cmap="coolwarm", annot=False)  
    plt.title("Correlation Heatmap")  
    plt.show()
```



1 Linear Regression -----

```
[17]: # pick numeric target for regression
linear_target = num_cols[0] if num_cols else None
```

```
[18]: print("Linear Regression Target:", linear_target)
```

Linear Regression Target: Avg VTAT

```
[19]: X = df[num_cols[1:]]
Y = df[linear_target]

model = LinearRegression()
model.fit(X, Y)

df['predicted_class'] = model.predict(X)
```

```

print(df)

coefficient = model.coef_[0]
intercept = model.intercept_
print("coefficient:", coefficient)
print("intercept:", intercept)

if X.shape[1] == 1:
    plt.scatter(X, Y, color='blue', label='Actual data')
    plt.plot(X, Y, color='red', linewidth=2, label=f'Linear regression: y = {coefficient:.2f}x + {intercept:.2f}')
else:
    feature_to_plot = X.columns[0]
    plt.scatter(X[feature_to_plot], Y, color='blue', label='Actual data')
    plt.plot(X[feature_to_plot], model.predict(X), color='red', linewidth=2, label='Linear regression')

plt.xlabel('Avg CTAT_Booking Value_Ride Distance_Driver Ratings_Customer Rating')
plt.ylabel(linear_target)
plt.title('Linear Regression Visualization')
plt.legend()
plt.grid(True)
plt.show()

```

| | Booking Status | Vehicle Type | Pickup Location | \ |
|--------|-----------------|---------------|------------------------|---|
| 0 | No Driver Found | eBike | Palam Vihar | |
| 1 | Incomplete | Go Sedan | Shastri Nagar | |
| 2 | Completed | Auto | Khandsa | |
| 3 | Completed | Premier Sedan | Central Secretariat | |
| 4 | Completed | Bike | Ghitorni Village | |
| ... | ... | ... | ... | |
| 149995 | Completed | Go Mini | MG Road | |
| 149996 | Completed | Go Mini | Golf Course Road | |
| 149997 | Completed | Go Sedan | Satguru Ram Singh Marg | |
| 149998 | Completed | Auto | Ghaziabad | |
| 149999 | Completed | Premier Sedan | Ashok Park Main | |

| | Drop Location | Avg VTAT | Avg CTAT | Booking Value | Ride Distance | \ |
|--------|-------------------|----------|----------|---------------|---------------|---|
| 0 | Jhilmil | 8.3 | 28.8 | 414.0 | 23.72 | |
| 1 | Gurgaon Sector 56 | 4.9 | 14.0 | 237.0 | 5.73 | |
| 2 | Malviya Nagar | 13.4 | 25.8 | 627.0 | 13.58 | |
| 3 | Inderlok | 13.1 | 28.5 | 416.0 | 34.02 | |
| 4 | Khan Market | 5.3 | 19.6 | 737.0 | 48.21 | |
| ... | ... | ... | ... | ... | ... | |
| 149995 | Ghitorni | 10.2 | 44.4 | 475.0 | 40.08 | |
| 149996 | Akshardham | 5.1 | 30.8 | 1093.0 | 21.31 | |


| | | | | | |
|--------|-------------------|-----|------|-------|-------|
| 149997 | Jor Bagh | 2.7 | 23.4 | 852.0 | 15.93 |
| 149998 | Saidulajab | 6.9 | 39.6 | 333.0 | 45.54 |
| 149999 | Gurgaon Sector 29 | 3.5 | 33.7 | 806.0 | 21.19 |

| | Driver Ratings | Customer Rating | Payment Method | predicted_class |
|--------|----------------|-----------------|----------------|-----------------|
| 0 | 4.3 | 4.5 | NaN | 8.431010 |
| 1 | 4.3 | 4.5 | UPI | 7.847132 |
| 2 | 4.9 | 4.9 | Debit Card | 8.149164 |
| 3 | 4.6 | 5.0 | UPI | 8.533886 |
| 4 | 4.1 | 4.3 | UPI | 8.569115 |
| ... | ... | ... | ... | ... |
| 149995 | 3.7 | 4.1 | Uber Wallet | 9.068923 |
| 149996 | 4.8 | 5.0 | UPI | 8.339190 |
| 149997 | 3.9 | 4.4 | Cash | 8.186736 |
| 149998 | 4.1 | 3.7 | UPI | 9.046387 |
| 149999 | 4.6 | 4.9 | Credit Card | 8.441368 |

[150000 rows x 12 columns]

coefficient: 0.022833522302575897

intercept: 7.8980921619235165



```
[20]: X = df['Avg CTAT'].values.reshape(-1, 1)
Y = df['Avg VTAT']

model = LinearRegression()
model.fit(X, Y)

#df['predicted_prob'] = model.predict_proba(X)[:,1]
df['predicted_class'] = model.predict(X)
print(df)

print("Intercept(B0):", model.intercept_)
print("Coefficient(B1):", model.coef_)

# Store the coefficient and intercept in variables before using them
coefficient = model.coef_[0] # Extract the coefficient value from the array
intercept = model.intercept_

plt.scatter(X, Y, color='blue', label='Actual data')
plt.plot(X, Y, color='red', linewidth=2, label=f'Linear regression: Y = {coefficient:.2f}X + {intercept:.2f}')

plt.xlabel('Avg CTAT')
plt.ylabel('Avg VTAT')
plt.title('Linear Regression Visualization')
plt.legend()
plt.grid(True)
plt.show()
```

| | Booking Status | Vehicle Type | Pickup Location | \ |
|--------|-------------------|---------------|------------------------|-------------------------------|
| 0 | No Driver Found | eBike | Palam Vihar | |
| 1 | Incomplete | Go Sedan | Shastri Nagar | |
| 2 | Completed | Auto | Khanda | |
| 3 | Completed | Premier Sedan | Central Secretariat | |
| 4 | Completed | Bike | Ghitorni Village | |
| ... | ... | ... | ... | |
| 149995 | Completed | Go Mini | MG Road | |
| 149996 | Completed | Go Mini | Golf Course Road | |
| 149997 | Completed | Go Sedan | Satguru Ram Singh Marg | |
| 149998 | Completed | Auto | Ghaziabad | |
| 149999 | Completed | Premier Sedan | Ashok Park Main | |
| | Drop Location | Avg VTAT | Avg CTAT | Booking Value Ride Distance \ |
| 0 | Jhilmil | 8.3 | 28.8 | 414.0 23.72 |
| 1 | Gurgaon Sector 56 | 4.9 | 14.0 | 237.0 5.73 |
| 2 | Malviya Nagar | 13.4 | 25.8 | 627.0 13.58 |
| 3 | Inderlok | 13.1 | 28.5 | 416.0 34.02 |
| 4 | Khan Market | 5.3 | 19.6 | 737.0 48.21 |


| | | | | | | |
|--------|-------------------|------|------|--------|-----|-------|
| ... | ... | ... | ... | ... | ... | ... |
| 149995 | Ghitorni | 10.2 | 44.4 | 475.0 | ... | 40.08 |
| 149996 | Akshardham | 5.1 | 30.8 | 1093.0 | ... | 21.31 |
| 149997 | Jor Bagh | 2.7 | 23.4 | 852.0 | ... | 15.93 |
| 149998 | Saidulajab | 6.9 | 39.6 | 333.0 | ... | 45.54 |
| 149999 | Gurgaon Sector 29 | 3.5 | 33.7 | 806.0 | ... | 21.19 |

| | Driver Ratings | Customer Rating | Payment Method | predicted_class |
|--------|----------------|-----------------|----------------|-----------------|
| 0 | 4.3 | 4.5 | NaN | 8.439408 |
| 1 | 4.3 | 4.5 | UPI | 8.065948 |
| 2 | 4.9 | 4.9 | Debit Card | 8.363707 |
| 3 | 4.6 | 5.0 | UPI | 8.431838 |
| 4 | 4.1 | 4.3 | UPI | 8.207257 |
| ... | ... | ... | ... | ... |
| 149995 | 3.7 | 4.1 | Uber Wallet | 8.833055 |
| 149996 | 4.8 | 5.0 | UPI | 8.489876 |
| 149997 | 3.9 | 4.4 | Cash | 8.303145 |
| 149998 | 4.1 | 3.7 | UPI | 8.711933 |
| 149999 | 4.6 | 4.9 | Credit Card | 8.563053 |

[150000 rows x 12 columns]

Intercept(B0): 7.71267486029648

Coefficient(B1): [0.02523379]



```
[21]: # Drop missing values for target
target = "Booking Value"
df_lr = df.dropna(subset=[target])

# Apply log transformation to target
df_lr["log_booking_value"] = np.log1p(df_lr[target]) # log(1+x) to handle 0s

# Features (drop target, status, and transformed column)
X = df_lr.drop(columns=["Booking Status", target, "log_booking_value"])
y = df_lr["log_booking_value"]
```

```
[22]: # Preprocessing
numeric_transformer = Pipeline([
    ("imputer", SimpleImputer(strategy="median")),
    ("scaler", StandardScaler())
])

categorical_transformer = Pipeline([
    ("imputer", SimpleImputer(strategy="most_frequent")),
    ("encoder", OneHotEncoder(handle_unknown="ignore"))
])

preprocessor = ColumnTransformer([
    ("num", numeric_transformer, [c for c in num_cols if c in X.columns]),
    ("cat", categorical_transformer, [c for c in cat_cols if c in X.columns])
])
```

```
[23]: # Train-test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
```

```
[24]: # Linear Regression pipeline
lin_reg_log_model = Pipeline([
    ("preprocessor", preprocessor),
    ("regressor", LinearRegression())
])
```

```
[25]: # Fit model
lin_reg_log_model.fit(X_train, y_train)
```

```
[25]: Pipeline(steps=[('preprocessor',
                      ColumnTransformer(transformers=[('num',
                                                       Pipeline(steps=[('imputer',
                                                               SimpleImputer(strategy='median')),
```

```

('scaler',
StandardScaler())]),

['Avg VTAT', 'Avg CTAT',
'Ride Distance',
'Driver Ratings',
'Customer Rating']),

('cat',
Pipeline(steps=[('imputer',
SimpleImputer(strategy='most_frequent')),

OneHotEncoder(handle_unknown='ignore')]]),

['Vehicle Type',
'Pickup Location',
'Drop Location',
'Payment Method']))),

('regressor', LinearRegression())))

```

[26]: # Predictions in log scale
`y_pred_log = lin_reg_log_model.predict(X_test)`


Convert back to original scale
`y_pred = np.expm1(y_pred_log) # reverse log1p`
`y_test_original = np.expm1(y_test)`

[27]: # Evaluation
`mse = mean_squared_error(y_test_original, y_pred)`
`rmse = np.sqrt(mse)`
`r2 = r2_score(y_test_original, y_pred)`

`print("Mean Squared Error:", mse)`
`print("Root Mean Squared Error:", rmse)`
`print("R2 Score:", r2)`

Mean Squared Error: 115599.10060619467
Root Mean Squared Error: 339.9986773594784
R² Score: -0.06743689382995588

[28]: # 1. Actual vs Predicted
`plt.figure(figsize=(7,7))`
`plt.scatter(y_test_original, y_pred, alpha=0.3)`
`plt.plot([y_test_original.min(), y_test_original.max()],
[y_test_original.min(), y_test_original.max()], 'r--')`
`plt.xlabel("Actual Booking Value")`
`plt.ylabel("Predicted Booking Value")`
`plt.title("Actual vs Predicted Booking Value (Log-Transformed Model)")`
`plt.show()`



2 1. Actual vs Predicted (Log-Transformed Model)

The predictions are heavily clustered in a narrow band (~300–500 range), while actual booking values span from very small amounts up to 4000+.

The red diagonal (perfect fit line) is far from most blue points → the model is underestimating high booking values and compressing predictions.


This suggests linear regression (even with log transformation) struggles to capture non-linear patterns in booking value.

```
[29]: # 2. Residuals Plot
residuals = y_test_original - y_pred
plt.figure(figsize=(7,5))
```

```

plt.scatter(y_pred, residuals, alpha=0.3)
plt.axhline(0, color='red', linestyle='--')
plt.xlabel("Predicted Booking Value")
plt.ylabel("Residuals")
plt.title("Residuals Plot (Log-Transformed Model)")
plt.show()

```



3 2. Residuals Plot

Residuals (errors) show a funnel-like shape: as predicted booking values increase, the spread of errors widens.

Many residuals are large positive values, meaning the model systematically under-predicts high fares.


Ideally, residuals should be randomly scattered around zero. Instead, we see structure → another sign linear regression is not flexible enough.

```
[30]: # 3. Residuals Distribution
plt.figure(figsize=(7,5))
sns.histplot(residuals, bins=50, kde=True)
```

```

plt.xlabel("Residual")
plt.title("Distribution of Residuals (Log-Transformed Model)")
plt.show()

```



4 3. Residuals Distribution

The residuals distribution is heavily right-skewed.

Peak around 0 (good), but a long positive tail indicates the model frequently misses high fares by a large margin.

This reinforces that linear regression is capturing the “average trend” but not the outliers and high-value bookings.

5 Key Takeaways

Log transformation helped stabilize variance a bit, but the model is still far from accurate.

The problem seems non-linear → Linear Regression isn't flexible enough.

High-value bookings are systematically underestimated.

Next step: try a non-linear regression model (e.g., Random Forest, Gradient Boosting, or XGBoost) which can capture complex interactions between features.

6 Logistic Regression

```
[31]: # Define target
df["target"] = df["Booking Status"].apply(lambda x: 1 if x == "Completed" else 0)

X = df.drop(columns=["Booking Status", "target"])
y = df["target"]
```

```
[32]: # Preprocessing pipelines
numeric_transformer = Pipeline([
    ("imputer", SimpleImputer(strategy="median")),
    ("scaler", StandardScaler())
])

categorical_transformer = Pipeline([
    ("imputer", SimpleImputer(strategy="most_frequent")),
    ("encoder", OneHotEncoder(handle_unknown="ignore"))
])

preprocessor = ColumnTransformer([
    ("num", numeric_transformer, [c for c in num_cols if c in X.columns]),
    ("cat", categorical_transformer, [c for c in cat_cols if c in X.columns])
])
```

```
[33]: X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

log_reg_model = Pipeline([
    ("preprocessor", preprocessor),
    ("classifier", LogisticRegression(max_iter=1000))
])

log_reg_model.fit(X_train, y_train)
```

```
[33]: Pipeline(steps=[('preprocessor',
                      ColumnTransformer(transformers=[('num',
                                                       Pipeline(steps=[('imputer',
                                                               SimpleImputer(strategy='median')),
                                                               ('scaler',
                                                               StandardScaler())])),
                                                       ['Avg VTAT', 'Avg CTAT'],
                                                       ]))]
```

```

        'Booking Value',
        'Ride Distance',
        'Driver Ratings',
        'Customer Rating']),
        ('cat',
        Pipeline(steps=[('imputer',
SimpleImputer(strategy='most_frequent')),
OneHotEncoder(handle_unknown='ignore')]),
        ['Vehicle Type',
        'Pickup Location',
        'Drop Location',
        'Payment Method'))),
        ('classifier', LogisticRegression(max_iter=1000))])

```

[34]: # Predictions

```
y_pred = log_reg_model.predict(X_test)
y_prob = log_reg_model.predict_proba(X_test)[:, 1]
```

[35]: # Evaluation

```
print("Accuracy:", accuracy_score(y_test, y_pred))
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

Accuracy: 0.8303666666666667

Confusion Matrix:

```
[[10307 1051]
 [ 4038 14604]]
```

Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.72 | 0.91 | 0.80 | 11358 |
| 1 | 0.93 | 0.78 | 0.85 | 18642 |
| accuracy | | | 0.83 | 30000 |
| macro avg | 0.83 | 0.85 | 0.83 | 30000 |
| weighted avg | 0.85 | 0.83 | 0.83 | 30000 |

[36]: # ROC-AUC curve (for binary classification)


```
if len(np.unique(y)) == 2:
    auc = roc_auc_score(y_test, y_prob)
    print("ROC-AUC:", auc)
    fpr, tpr, _ = roc_curve(y_test, y_prob)
    plt.plot(fpr, tpr, label=f"ROC Curve (AUC={auc:.2f})")
```

```

plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend()
plt.show()

```

ROC-AUC: 0.8220642347949073



7 KMeans Clustering -----

```


[37]: # Select features for clustering (exclude IDs, targets, non-numeric)
features = df.drop(columns=["Booking ID", "Booking Value"], errors="ignore") # ↵
# drop ID & target
X = features.select_dtypes(include=['int64', 'float64'])

[38]: # Standardize data
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

```

```
[39]: # Elbow method to find optimal k
inertia = []
K = range(1, 11)
for k in K:
    km = KMeans(n_clusters=k, random_state=42, n_init=10)
    km.fit(X_scaled)
    inertia.append(km.inertia_)
```

```
[40]: plt.figure(figsize=(6,4))
plt.plot(K, inertia, marker="o")
plt.xlabel("Number of Clusters (k)")
plt.ylabel("Inertia (SSE)")
plt.title("Elbow Method for Optimal k")
plt.show()
```



```
[41]: # Fit KMeans with chosen k (say k=3 or 4 after elbow plot)
kmeans = KMeans(n_clusters=4, random_state=42, n_init=10)
df["Cluster"] = kmeans.fit_predict(X_scaled)
```

```
[42]: # Reduce dimensions for visualization
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_scaled)
```

```
[43]: # Example: interpret clusters by their average Booking Value
cluster_summary = df.groupby("Cluster")["Booking Value"].mean().sort_values()
print(cluster_summary)

# Map cluster numbers → meaningful names
cluster_labels = {
    cluster_summary.index[0]: "Low Fare Riders",
    cluster_summary.index[1]: "Mid Fare Riders",
    cluster_summary.index[2]: "High Fare Riders",
    cluster_summary.index[3]: "Premium Riders"
}

df["Cluster Name"] = df["Cluster"].map(cluster_labels)

plt.figure(figsize=(7,5))
for cluster, name in cluster_labels.items():
    plt.scatter(
        X_pca[df["Cluster"]==cluster,0],
        X_pca[df["Cluster"]==cluster,1],
        label=name, alpha=0.6
    )
plt.scatter(X_pca[:,0], X_pca[:,1], c=df["Cluster"], cmap="viridis", alpha=0.6)
plt.xlabel("PCA Component 1")
plt.ylabel("PCA Component 2")
plt.title("KMeans Clusters of Ride Bookings")
plt.legend(title="Cluster Type")
plt.show()
```

```
Cluster
1    424.300054
2    507.515531
3    508.015164
0    509.015239
Name: Booking Value, dtype: float64

C:\Users\user\anaconda3\Lib\site-packages\IPython\core\pylabtools.py:170:
UserWarning: Creating legend with loc="best" can be slow with large amounts of
data.
    fig.canvas.print_figure(bytes_io, **kw)
```



8 Insights from the Graph

1. Clear Segmentation of Riders

The clusters are well-separated, meaning the KMeans algorithm has successfully grouped riders based on their booking characteristics (likely influenced by fare, ride distance, or duration).

Each cluster shows a distinct behavioral pattern, useful for targeted strategies.

2. Low Fare Riders (Blue Cluster, far left)

This group is clearly isolated, suggesting they have consistently low booking values.

Likely represent budget-conscious commuters (e.g., daily short-distance riders).

3. Mid Fare Riders (Yellow Cluster, center)

Positioned between low and high fare groups.

This group has moderate booking values, possibly occasional riders or medium-distance commuters.

4. High Fare Riders (Green Cluster, right side)

Spread across the upper-right side, showing higher variability in ride behavior.

Indicates frequent travelers with higher-than-average spending.

5. Premium Riders (Red/Purple Cluster, overlapping with High Fare)

Found on the far-right edge, they represent the highest booking values.

Likely include luxury service users, long-distance travelers, or corporate accounts.

Though close to high fare riders, they form their own elite cluster.

6. Business Implications

Low Fare Riders → Can be offered loyalty rewards or discounts to encourage frequent use.

Mid Fare Riders → Upselling opportunities (e.g., premium rides, bundled packages).

High Fare Riders → Personalized offers for long-distance trips.

Premium Riders → Exclusive benefits (priority booking, premium customer support).

9 Random Forest Classifier

```
[44]: # Make a copy
df_encoded = df.copy()

# Encode categorical variables
for col in df_encoded.select_dtypes(include=['object']).columns:
    df_encoded[col] = LabelEncoder().fit_transform(df_encoded[col].astype(str))

# Define Features & Target
X = df_encoded.drop(columns=["Booking Status"])      # features
y = df_encoded["Booking Status"]                      # target

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,random_state=42)

# Model
rf_clf = RandomForestClassifier(n_estimators=100, random_state=42)
rf_clf.fit(X_train, y_train)

# Predictions
y_pred = rf_clf.predict(X_test)

# Evaluation
print("Accuracy:", accuracy_score(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred))

# Feature Importance
importances = rf_clf.feature_importances_
```

```

indices = np.argsort(importances)[::-1]
plt.figure(figsize=(10,5))
plt.title("Feature Importances")
plt.bar(range(X.shape[1]), importances[indices])
plt.xticks(range(X.shape[1]), X.columns[indices], rotation=90)
plt.xlabel("Features", fontsize=12)
plt.ylabel("Importance Score", fontsize=12)
plt.show()

```

Accuracy: 0.9621666666666666

Classification Report:


| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.92 | 0.55 | 0.68 | 2077 |
| 1 | 0.85 | 0.97 | 0.90 | 5431 |
| 2 | 1.00 | 1.00 | 1.00 | 18642 |
| 3 | 1.00 | 1.00 | 1.00 | 1752 |
| 4 | 0.95 | 1.00 | 0.97 | 2098 |
| accuracy | | | 0.96 | 30000 |
| macro avg | 0.94 | 0.90 | 0.91 | 30000 |
| weighted avg | 0.96 | 0.96 | 0.96 | 30000 |

Confusion Matrix:

```

[[ 1135   923     0     0    19]
 [ 104   5243     0     0    84]
 [  0     0  18642     0     0]
 [  0     0     0  1752     0]
 [  1     4     0     0 2093]]

```



10 Random Forest Regressor ——

```
[46]: # --- 1) Prep data ---
df_rf = df.copy()

# Drop rows with missing target
df_rf = df_rf.dropna(subset=["Booking Value"])

# Optional: drop pure ID columns that don't help prediction
for col in ["Booking ID", "Customer ID"]:
    if col in df_rf.columns:
        df_rf.drop(columns=[col], inplace=True)

# Impute missing values BEFORE encoding
for col in df_rf.columns:
    if df_rf[col].dtype == "object":
        df_rf[col] = df_rf[col].fillna("Unknown")
    else:
        df_rf[col] = df_rf[col].fillna(df_rf[col].median())

# Label-encode ALL object columns (trees handle integer-coded categories well)
df_enc = df_rf.copy()
```

```

for col in df_enc.select_dtypes(include=["object"]).columns:
    df_enc[col] = LabelEncoder().fit_transform(df_enc[col].astype(str))

# --- 2) Features & target ---
X = df_enc.drop(columns=["Booking Value"])
y = df_enc["Booking Value"]

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

# --- 3) Model ---
rf_reg = RandomForestRegressor(
    n_estimators=200,
    max_depth=None,
    random_state=42,
    n_jobs=-1
)
rf_reg.fit(X_train, y_train)

# --- 4) Predict & evaluate (compat with older sklearn) ---
y_pred = rf_reg.predict(X_test)

mse = mean_squared_error(y_test, y_pred)      # no 'squared' arg in older
                                              # versions
rmse = np.sqrt(mse)                          # compute RMSE manually
mae = mean_absolute_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print("R² Score:", round(r2, 4))
print("MSE:", round(mse, 2))
print("RMSE:", round(rmse, 2))
print("MAE:", round(mae, 2))

# --- 5) Feature importance plot (with axis labels) ---
importances = rf_reg.feature_importances_
indices = np.argsort(importances)[::-1]

top_n = min(20, X.shape[1])  # show top 20 for readability
top_idx = indices[:top_n]

plt.figure(figsize=(10,5))
plt.title("Top Feature Importances - Random Forest Regressor", fontsize=14)
plt.bar(range(top_n), importances[top_idx])
plt.xticks(range(top_n), X.columns[top_idx], rotation=90)
plt.xlabel("Features", fontsize=12)
plt.ylabel("Importance Score", fontsize=12)

```

```

plt.tight_layout()
plt.show()

# --- 6) Predicted vs Actual scatter (helpful sanity check) ---
plt.figure(figsize=(6,6))
plt.scatter(y_test, y_pred, alpha=0.3)
m = min(y_test.min(), y_pred.min())
M = max(y_test.max(), y_pred.max())
plt.plot([m, M], [m, M], 'r--')
plt.xlabel("Actual Booking Value")
plt.ylabel("Predicted Booking Value")
plt.title("Actual vs Predicted - Random Forest Regressor")
plt.tight_layout()
plt.show()


```

R² Score: -0.0116

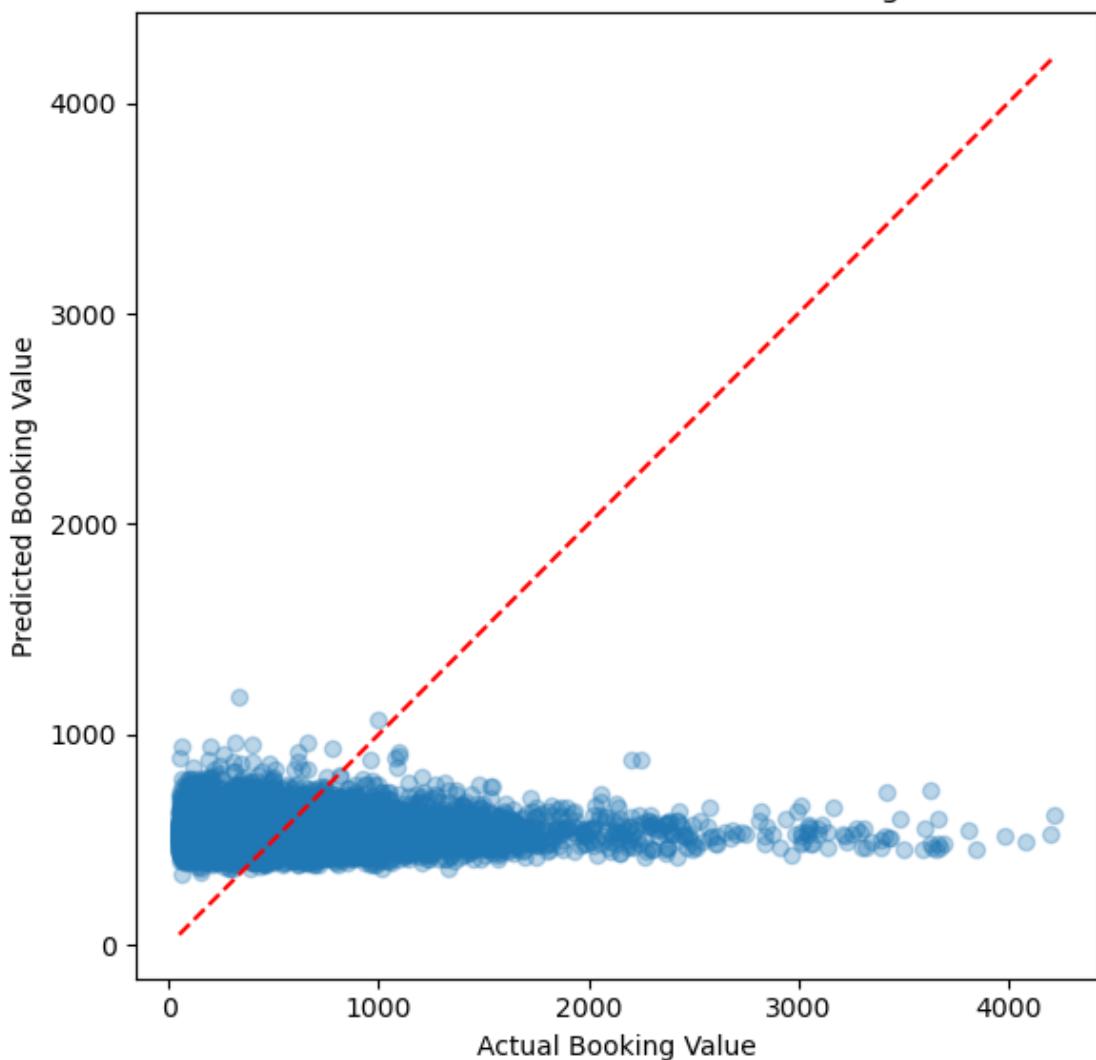
MSE: 109546.85

RMSE: 330.98

MAE: 200.8



Actual vs Predicted — Random Forest Regressor



[]: